

Realistic example

So far, you have seen the basics of manipulating data frames with our cat data; now let's use those skills to digest a more realistic dataset. Let's read in the `gapminder` dataset that we downloaded previously:

```
gapminder <- read.csv("data/gapminder_data.csv", stringsAsFactors = TRUE)
```

Miscellaneous Tips

- Another type of file you might encounter are tab-separated value files (.tsv). To specify a tab as a separator, use `"\\t"` or `read.delim()`.
- Files can also be downloaded directly from the Internet into a local folder of your choice onto your computer using the `download.file` function. The `read.csv` function can then be executed to read the downloaded file from the download location, for example,

```
download.file("https://raw.githubusercontent.com/swcarpentry/r-novice-gapminder/gh-pages/_episodes_rmd/data/gapminder_data.csv", destfile = "data/gapminder_data.csv")
gapminder <- read.csv("data/gapminder_data.csv", stringsAsFactors = TRUE)
```

- Alternatively, you can also read in files directly into R from the Internet by replacing the file paths with a web address in `read.csv`. One should note that in doing this no local copy of the csv file is first saved onto your computer. For example,

```
gapminder <- read.csv("https://raw.githubusercontent.com/swcarpentry/r-novice-gapminder/gh-pages/_episodes_rmd/data/gapminder_data.csv", stringsAsFactors = TRUE)
```

- You can read directly from excel spreadsheets without converting them to plain text first by using the [readxl](#) package.

Let's investigate `gapminder` a bit; the first thing we should always do is check out what the data looks like with `str`:

```
str(gapminder)
```

```
'data.frame': 1704 obs. of 6 variables:
 $ country : Factor w/ 142 levels "Afghanistan",...: 1 1 1 1 1 1 1 1 1 ...
 $ year : int 1952 1957 1962 1967 1972 1977 1982 1987 1992 1997 ...
 $ pop : num 8425333 9240934 10267083 11537966 13079460 ...
 $ continent: Factor w/ 5 levels "Africa","Americas",...: 3 3 3 3 3 3 3 3 3 ...
 $ lifeExp : num 28.8 30.3 32 34 36.1 ...
 $ gdpPercap: num 779 821 853 836 740 ...
```

An additional method for examining the structure of `gapminder` is to use the `summary` function. This function can be used on various objects in R. For data frames, `summary` yields a numeric, tabular, or descriptive summary of each column. Factor columns are summarized by the number of items in each level, numeric or integer columns by the descriptive statistics (quartiles and mean), and character columns by its length, class, and mode.

```
summary(gapminder$country)
```

Afghanistan	Albania	Algeria
12	12	12
Angola	Argentina	Australia
12	12	12
Austria	Bahrain	Bangladesh
12	12	12
Belgium	Benin	Bolivia
12	12	12
Bosnia and Herzegovina	Botswana	Brazil
12	12	12
Bulgaria	Burkina Faso	Burundi
12	12	12
Cambodia	Cameroon	Canada
12	12	12
Central African Republic	Chad	Chile
12	12	12
China	Colombia	Comoros
12	12	12
Congo Dem. Rep.	Congo Rep.	Costa Rica
12	12	12
Cote d'Ivoire	Croatia	Cuba
12	12	12
Czech Republic	Denmark	Djibouti
12	12	12
Dominican Republic	Ecuador	Egypt
12	12	12
El Salvador	Equatorial Guinea	Eritrea
12	12	12
Ethiopia	Finland	France
12	12	12
Gabon	Gambia	Germany
12	12	12
Ghana	Greece	Guatemala
12	12	12
Guinea	Guinea-Bissau	Haiti
12	12	12
Honduras	Hong Kong China	Hungary
12	12	12
Iceland	India	Indonesia
12	12	12
Iran	Iraq	Ireland
12	12	12
Israel	Italy	Jamaica
12	12	12
Japan	Jordan	Kenya
12	12	12
Korea Dem. Rep.	Korea Rep.	Kuwait
12	12	12
Lebanon	Lesotho	Liberia
12	12	12
Libya	Madagascar	Malawi
12	12	12
Malaysia	Mali	Mauritania
12	12	12
Mauritius	Mexico	Mongolia
12	12	12
Montenegro	Morocco	Mozambique
12	12	12
Myanmar	Namibia	Nepal
12	12	12
Netherlands	New Zealand	Nicaragua
12	12	12
Niger	Nigeria	Norway
12	12	12
Oman	Pakistan	Panama
12	12	12
(Other)		
516		

Along with the `str` and `summary` functions, we can examine individual columns of the data frame with our `typeof` function:

```
typeof(gapminder$year)
```

```
[1] "integer"
```

```
typeof(gapminder$country)
```

```
[1] "integer"
```

```
str(gapminder$country)
```

```
Factor w/ 142 levels "Afghanistan",...: 1 1 1 1 1 1 1 1 1 1 ...
```

We can also interrogate the data frame for information about its dimensions; remembering that `str(gapminder)` said there were 1704 observations of 6 variables in `gapminder`, what do you think the following will produce, and why?

```
length(gapminder)
```

```
[1] 6
```

A fair guess would have been to say that the length of a data frame would be the number of rows it has (1704), but this is not the case; remember, a data frame is a *list of vectors and factors*:

```
typeof(gapminder)
```

```
[1] "list"
```

When `length` gave us 6, it's because `gapminder` is built out of a list of 6 columns. To get the number of rows and columns in our dataset, try:

```
nrow(gapminder)
```

```
[1] 1704
```

```
ncol(gapminder)
```

```
[1] 6
```

Or, both at once:

```
dim(gapminder)
```

```
[1] 1704    6
```

We'll also likely want to know what the titles of all the columns are, so we can ask for them later:

```
colnames(gapminder)
```

```
[1] "country" "year"    "pop"     "continent" "lifeExp"  "gdpPercap"
```

At this stage, it's important to ask ourselves if the structure R is reporting matches our intuition or expectations; do the basic data types reported for each column make sense? If not, we need to sort any problems out now before they turn into bad surprises down the road, using what we've learned about how R interprets data, and the importance of *strict consistency* in how we record our data.

Once we're happy that the data types and structures seem reasonable, it's time to start digging into our data proper. Check out the first few lines:

```
head(gapminder)
```

	country	year	pop	continent	lifeExp	gdpPercap
1	Afghanistan	1952	8425333	Asia	28.801	779.4453
2	Afghanistan	1957	9240934	Asia	30.332	820.8530
3	Afghanistan	1962	10267083	Asia	31.997	853.1007
4	Afghanistan	1967	11537966	Asia	34.020	836.1971
5	Afghanistan	1972	13079460	Asia	36.088	739.9811
6	Afghanistan	1977	14880372	Asia	38.438	786.1134

Challenge 3

It's good practice to also check the last few lines of your data and some in the middle. How would you do this?

Searching for ones specifically in the middle isn't too hard, but we could ask for a few lines at random. How would you code this?

Solution to Challenge 3

To check the last few lines it's relatively simple as R already has a function for this:

```
tail(gapminder)
tail(gapminder, n = 15)
```

What about a few arbitrary rows just in case something is odd in the middle?

Tip: There are several ways to achieve this.

The solution here presents one form of using nested functions, i.e. a function passed as an argument to another function. This might sound like a new concept, but you are already using it! Remember `my_dataframe[rows, cols]` will print to screen your data frame with the number of rows and columns you asked for (although you might have asked for a range or named columns for example). How would you get the last row if you don't know how many rows your data frame has? R has a function for this. What about getting a (pseudorandom) sample? R also has a function for this.

```
gapminder[sample(nrow(gapminder), 5), ]
```